

Refresher on R essentials

PSCI 2301: Quantitative Political Science II

Prof. Brenton Kenkel

brenton.kenkel@gmail.com

Vanderbilt University

January 8, 2025

Today's agenda

1. Getting data into R
2. Manipulating and cleaning data with tidyverse
3. Data visualization with ggplot

Before we go further: Packages

All R code I write in this course will assume you've got the **tidyverse** R package installed and loaded.

```
library("tidyverse")
```

If that gives you an error message like...

```
Error in library("tidyverse") : there is no package called 'tidyverse'
```

...then you need to install the package by running

```
install.packages("tidyverse")
```

Getting data into R

Data from the public Internet

If you have a direct link to a CSV file, you can plug it into `read_csv()`.

Remember to put the URL inside quote marks.

```
df_lottery <- read_csv("https://data.ny.gov/api/views/5xaw-6ayf/rows.csv")
print(df_lottery)
```

```
# A tibble: 2,360 × 4
  `Draw Date` `Winning Numbers` `Mega Ball` Multiplier
  <chr>       <chr>             <chr>      <chr>
1 09/25/2020 20 36 37 48 67      16         02
2 09/29/2020 14 39 43 44 67      19         03
3 10/02/2020 09 38 47 49 68      25         02
4 10/06/2020 15 16 18 39 59      17         03
5 10/09/2020 05 11 25 27 64      13         02
# i 2,355 more rows
```

Data from a downloaded file

Often you can't `read_csv()` directly from a URL:

- Files behind logins, e.g. on the course Brightspace
- Files within a zipped directory
- Files that somebody emailed you

In this situation, you need to be mindful of your **working directory**

Use `getwd()` to find out where R is looking for files, `setwd()` to change it

```
getwd()
```

```
[1] "/home/brenton/Dropbox/courses/qps2/slides/01_02_r_refresher"
```

Data from a downloaded file: Practice

1. Create a directory somewhere you can find it
2. Download the `anes2020.csv` file from Brightspace and put it there
3. Set as R's working directory
 - Windows:
`setwd('C:/path/to/directory')`
 - Mac:
`setwd('~ /path/to/directory')`
 - ... or just navigate in RStudio and set it that way
4. Check that `df_anes <- read_csv("anes2020.csv")` works

```
df_anes <- read_csv("anes2020.csv")
```

```
print(df_anes)
```

```
# A tibble: 8,280 × 32
   id state      female  lgbt race    age
  <dbl> <chr>      <dbl> <dbl> <chr> <dbl>
1     1 Oklahoma        0     0 Hisp...  46
2     2 Idaho          1     0 Asian   37
3     3 Virginia        1     0 White  40
4     4 Californ...    0     0 Asian  41
5     5 Colorado        0     0 Nati... 72
# i 8,275 more rows
# i 26 more variables: education <chr>,
#   employed <dbl>, hours_worked <dbl>, ...
```

`read_csv` **versus** `read.csv`

I always use tidyverse's `read_csv` (w/ underscore) instead of R's built-in `read.csv` (w/ period)

- Automatically stores data frame as “tibble” \rightsquigarrow better output display
- Does *not* automatically encode text as “factor”
- Works faster + shows progress bar for large datasets

Manipulating data

Basics of data manipulation

Use `$` to extract a single column

```
df_anes$age
```

```
[1] 46 37 40 41 72 71 37 45 70 43 37 55 30 38 41 66 54 55 62 80 31 80 24 55 59  
[ reached getOption("max.print") -- omitted 8255 entries ]
```

Use square brackets `[]` to extract individual value(s)

```
df_anes$age[5]      # age of the 5'th row of the data
```

```
[1] 72
```

```
df_anes$age[1:10]   # first 10 ages in the data
```

```
[1] 46 37 40 41 72 71 37 45 70 43
```

Useful data summaries

```
mean(df_anes$age, na.rm = TRUE) # average/mean
```

```
[1] 51.58522
```

```
median(df_anes$age, na.rm = TRUE) # median
```

```
[1] 52
```

```
sd(df_anes$age, na.rm = TRUE) # standard deviation
```

```
[1] 17.20718
```

```
table(df_anes$race) # counts of values -> can also do w/ summarize() or count()
```

Asian	Black	Hispanic	Multiracial	Native American
284	726	762	271	172
White				
5963				

Reducing data by row or column

```
filter(df_anes, age >= 75) # by row
```

```
# A tibble: 793 × 32
```

	id	state	female	lgbt	race	age	education	employed	hours_worked	watch_tucker
	<dbl>	<chr>	<dbl>	<dbl>	<chr>	<dbl>	<chr>	<dbl>	<dbl>	<dbl>
1	20	Wisconsin	1	0	White	80	High sch...	0	0	0
2	22	California	1	0	Hispanic	80	Less tha...	0	0	0
3	47	Pennsylvania	0	0	White	79	Some col...	0	0	0
4	63	Tennessee	0	0	White	80	Graduate...	0	0	0
5	68	California	1	0	White	78	Graduate...	0	0	0

```
# i 788 more rows
```

```
# i 22 more variables: watch_maddow <dbl>, therm_biden <dbl>, therm_trump <dbl>, ...
```

```
select(df_anes, female, age, education) # by column
```

```
# A tibble: 8,280 × 3
```

	female	age	education
	<dbl>	<dbl>	<chr>
1	0	46	Bachelor's degree
2	1	37	Some college
3	1	40	High school
4	0	41	Some college
5	0	72	Graduate degree

```
# i 8,275 more rows
```

Chaining commands with the pipe

```
df_anes |>
  filter(age >= 75) |>
  select(female, age, education)
```

```
# A tibble: 793 × 3
  female   age education
  <dbl> <dbl> <chr>
1       1    80 High school
2       1    80 Less than high school
3       0    79 Some college
4       0    80 Graduate degree
5       1    78 Graduate degree
# i 788 more rows
```

Different pipes

I use R's built-in pipe `|>`. Online you'll find a lot of code using the tidyverse pipe `%>%`. Both are fine and do essentially the same thing.

Changing and adding columns

```
df_anes |>
  mutate(female = if_else(female == 1, "yes", "no"),
         age_in_days = age * 365,
         employment_type = case_when(
           hours_worked == 0 ~ "unemployed",
           hours_worked < 32 ~ "part-time",
           hours_worked >= 32 ~ "full-time"
         )) |>
  relocate(female, age_in_days, employment_type) # put these cols first
```

A tibble: 8,280 × 34

	female	age_in_days	employment_type	id	state	lgbt	race	age	education	employed
	<chr>	<dbl>	<chr>	<dbl>	<chr>	<dbl>	<chr>	<dbl>	<chr>	<dbl>
1	no	16790	full-time	1	Oklahoma	0	Hispanic	46	Bachelor...	1
2	yes	13505	full-time	2	Idaho	0	Asian	37	Some col...	1
3	yes	14600	unemployed	3	Virginia	0	White	40	High sch...	0
4	no	14965	full-time	4	California	0	Asian	41	Some col...	1
5	no	26280	unemployed	5	Colorado	0	Native ...	72	Graduate...	0

i 8,275 more rows

i 24 more variables: hours_worked <dbl>, watch_tucker <dbl>, watch_madow <dbl>, ...

Summaries by group

```
df_anes |>
  group_by(race) |>
  summarize(n_respondents = n(),
            avg_trump_feeling = mean(therm_trump, na.rm = TRUE),
            sd_trump_feeling = sd(therm_trump, na.rm = TRUE))
```

A tibble: 7 × 4

	race	n_respondents	avg_trump_feeling	sd_trump_feeling
	<chr>	<int>	<dbl>	<dbl>
1	Asian	284	34.0	36.1
2	Black	726	15.0	25.2
3	Hispanic	762	32.2	36.7
4	Multiracial	271	34.8	38.8
5	Native American	172	42.2	38.1
6	White	5963	45.0	41.2
7	<NA>	102	41.3	38.3

Making changes stick

R commands almost never change a data frame in memory. The results of `filter()`, `select()`, `mutate()`, etc., will disappear unless you use `<-` to overwrite the original data frame or create a new one.

```
df_anes_women <- df_anes |>
  filter(female == 1)
```

df_anes

```
# A tibble: 8,280 × 32
   id state    female  lgbt race    age
  <dbl> <chr>    <dbl> <dbl> <chr>  <dbl>
1     1 Oklahoma      0     0 Hispanic  46
2     2 Idaho        1     0 Asian    37
3     3 Virginia     1     0 White    40
4     4 California   0     0 Asian    41
5     5 Colorado     0     0 Native ...  72
# i 8,275 more rows
# i 26 more variables: education <chr>,
#   employed <dbl>, hours_worked <dbl>, ...
```

df_anes_women

```
# A tibble: 4,450 × 32
   id state    female  lgbt race    age
  <dbl> <chr>    <dbl> <dbl> <chr>  <dbl>
1     2 Idaho      1     0 Asian    37
2     3 Virginia   1     0 White    40
3     6 Texas      1     0 White    71
4     7 Wisconsin  1     0 White    37
5     8 <NA>      1     0 White    45
# i 4,445 more rows
# i 26 more variables: education <chr>,
#   employed <dbl>, hours_worked <dbl>, ...
```


Other helpful data manipulation commands

- `case_match()` to code one column based on values of another
- `arrange()` for reordering rows
- `pivot_wider()` and `pivot_longer()` for reshaping data frames
- `left_join()` for merging data frames
- `group_by() |> mutate()` to add columns based on group-level calculations



Additional info on these commands

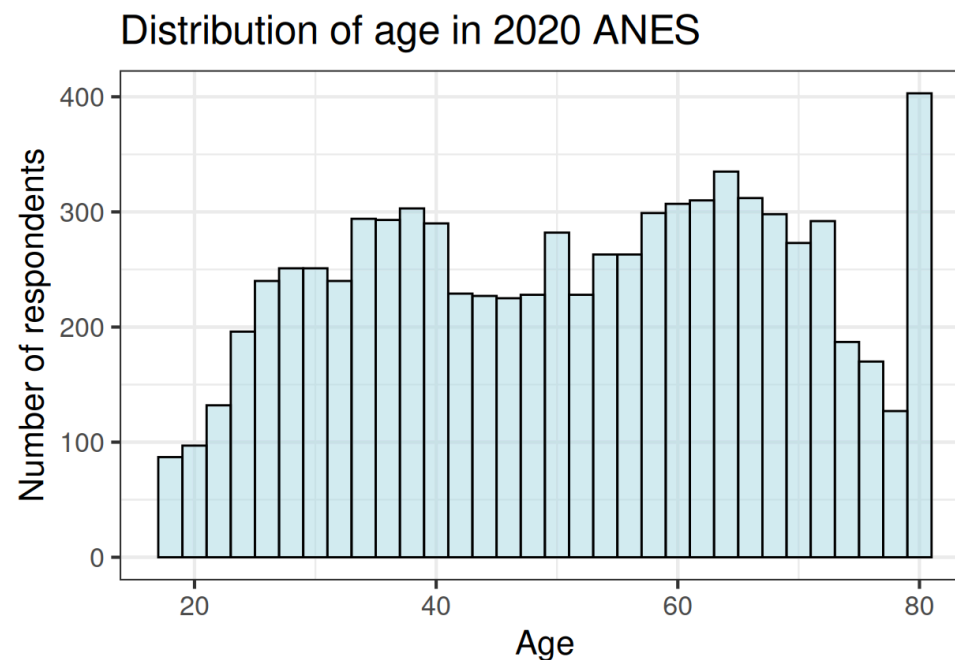
- The “Data Wrangling” notes from PSCI 2300 posted to Brightspace
- Lecture notes from [my graduate stats class](#)
- “Data Transformation” and “Data Tidying” chapters of [R for Data Science](#)

Data visualization

Visualizing a single variable

Continuous variables: Histogram

```
ggplot(df_anes, aes(x = age)) +  
  geom_histogram(color = "black",  
                 fill = "lightblue",  
                 alpha = 0.5,  
                 binwidth = 2) +  
  labs(x = "Age",  
        y = "Number of respondents",  
        title = "Distribution of age in 2020 ANES")
```



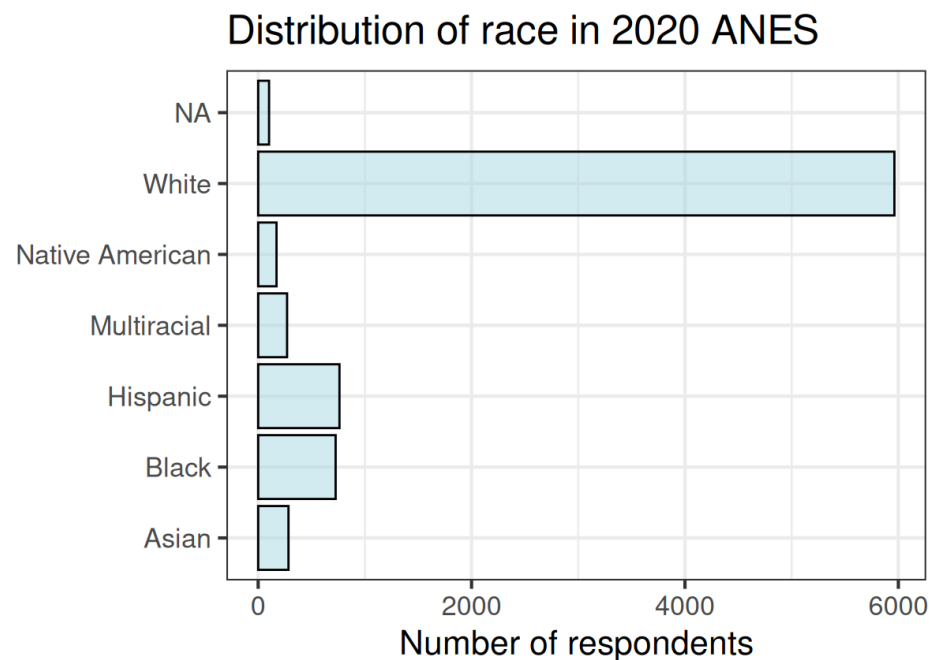
⚠ ggplot syntax

ggplot commands are separated by addition `+`, not the pipe `|>`.

Visualizing a single variable

Categorical variables: Bar chart


```
ggplot(df_anes, aes(y = race)) +  
  geom_bar(color = "black",  
           fill = "lightblue",  
           alpha = 0.5) +  
  labs(x = "Number of respondents",  
       y = "",  
       title = "Distribution of race in 2020 ANES")
```



Visualizing relationships

Two continuous variables: Scatterplot


```
1 ggplot(df_anes,  
2       aes(x = therm_biden, y = therm_trump))  
3 geom_point() +  
4 labs(x = "Feelings toward Biden",  
5       y = "Feelings toward Trump",  
6       title = "Shocker: Liking Biden predicts
```



Visualizing relationships

Two continuous variables: Scatterplot

```
1 ggplot(df_anes,  
2       aes(x = therm_biden, y = therm_trump))  
3   geom_point(position = "jitter",  
4             alpha = 0.1) +  
5   labs(x = "Feelings toward Biden",  
6        y = "Feelings toward Trump",  
7        title = "Shocker: Liking Biden predicts
```



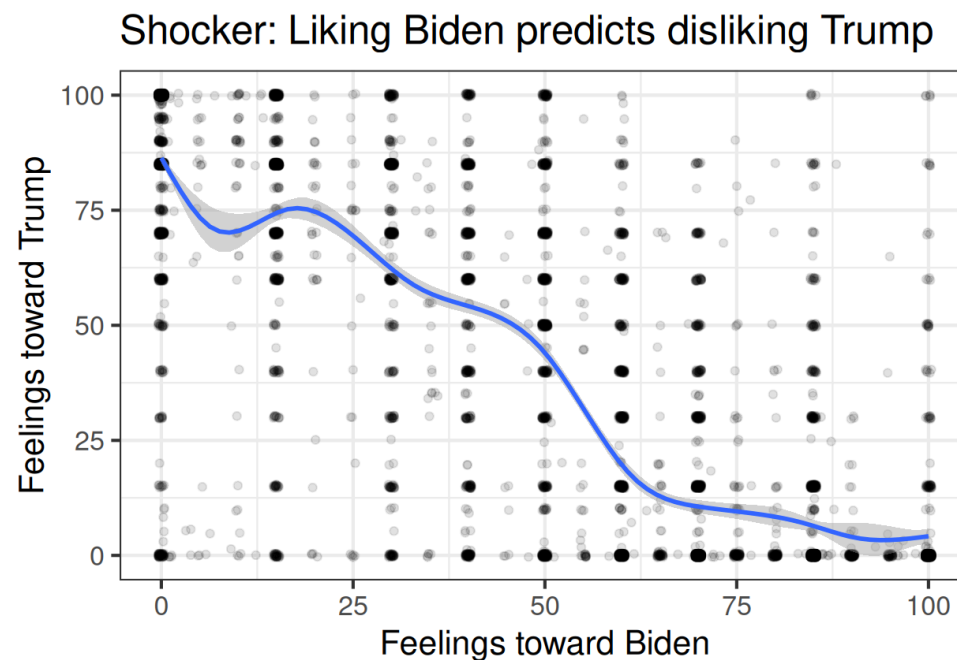
i Jitter and transparency

When data is clumpy with lots of overlapping values, jitter the point locations and/or make points semi-transparent to see relationships better.

Visualizing relationships

Two continuous variables: Scatterplot

```
1 ggplot(df_anes,  
2       aes(x = therm_biden, y = therm_trump))  
3   geom_point(position = "jitter",  
4             alpha = 0.1) +  
5   geom_smooth() +  
6   labs(x = "Feelings toward Biden",  
7        y = "Feelings toward Trump",  
8        title = "Shocker: Liking Biden predicts
```




i Smoothing lines

Use `geom_smooth()` for a flexible trend line, or `geom_smooth(method = "lm")` for the linear regression line.

Visualizing relationships

Continuous and categorical variable: Box plot


```
1 ggplot(df_anes, aes(x = hours_worked, y = educ
2   geom_boxplot() +
3   labs(x = "Hours worked per week",
4     y = "Highest educational attainment",
5     title = "Who works the longest hours?")
```



Visualizing relationships

Continuous and categorical variable: Box plot

```
1 ggplot(df_anes, aes(x = hours_worked, y = educ
2   geom_boxplot() +
3   scale_y_discrete(limits = c("Less than high
4     "High school",
5     "Some college",
6     "Bachelor's degr
7     "Graduate degree
8   labs(x = "Hours worked per week",
9     y = "Highest educational attainment",
10    title = "Who works the longest hours?"))
```




i Reordering a categorical variable

Use `scale_x_discrete(limits = c(...))` or `scale_y_discrete(limits = c(...))` to change the order the categories appear in.

Visualizing relationships

Continuous and categorical variable: Faceted histogram


```
1 ggplot(df_anes, aes(x = hours_worked)) +  
2   geom_histogram(binwidth = 5) +  
3   facet_wrap(~ education,  
4             scales = "free",  
5             ncol = 2) +  
6   labs(x = "Hours worked per week",  
7        y = "Number of respondents",  
8        title = "Distribution of working hours
```



Visualizing relationships

Continuous and categorical variable: Faceted histogram

```
1 ggplot(df_anes, aes(x = hours_worked)) +  
2   geom_histogram(binwidth = 5) +  
3   facet_wrap(~ fct_relevel(education,  
4     "Less than high sch  
5     "High school",  
6     "Some college",  
7     "Bachelor's degree"  
8     scales = "free",  
9     ncol = 2) +  
10  labs(x = "Hours worked per week",  
11       y = "Number of respondents",  
12       title = "Distribution of working hours
```



i Reordering facets

A bit fussier than reordering categories on an axis—use `fct_relevel()` within the call to `facet_wrap()`.

Visualizing relationships

Continuous and categorical variable: Bar chart summary

```
1 df_anes |>
2   group_by(education) |>
3   summarize(avg_hours = mean(hours_worked,
4                               na.rm = TRUE))
```

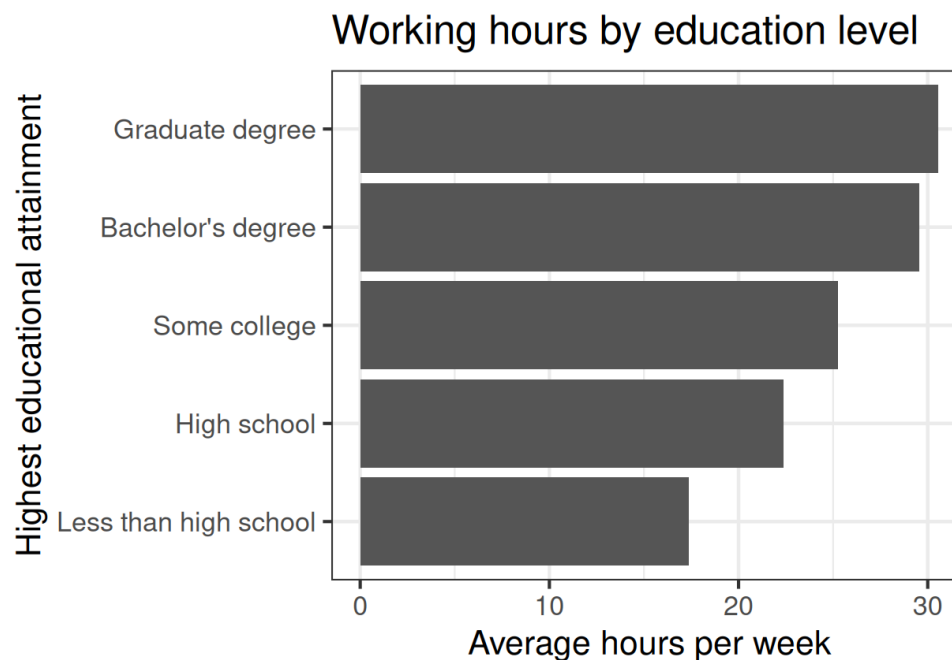
A tibble: 6 × 2

education	avg_hours
<chr>	<dbl>
1 Bachelor's degree	29.6
2 Graduate degree	30.6
3 High school	22.4
4 Less than high school	17.4
5 Some college	25.2
6 <NA>	15.9

Visualizing relationships

Continuous and categorical variable: Bar chart summary

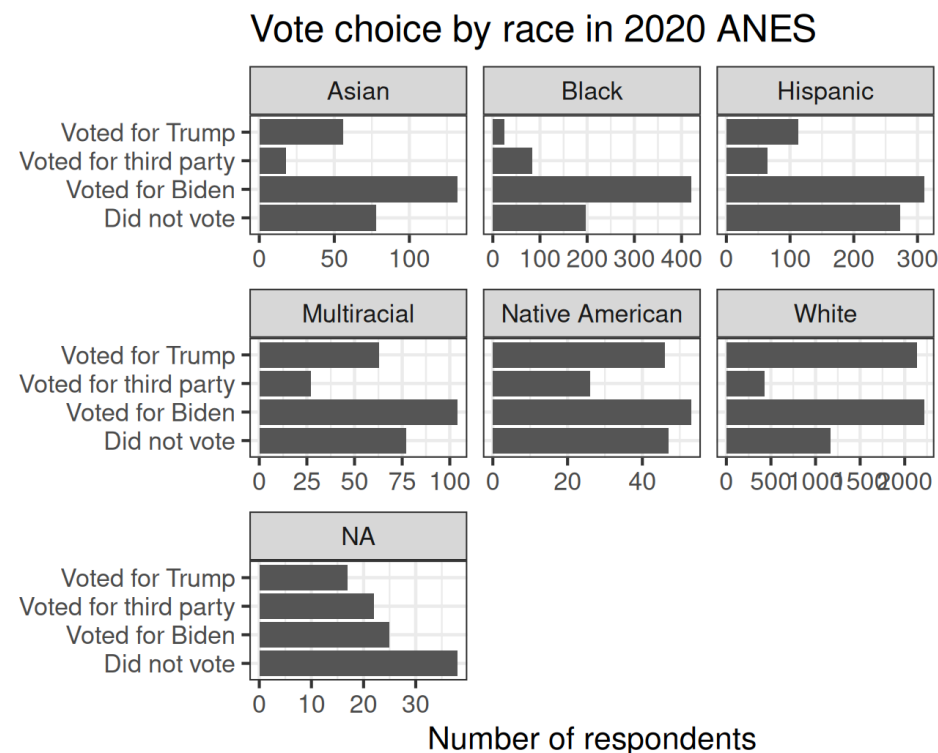
```
1 df_anes |>
2   group_by(education) |>
3   summarize(avg_hours = mean(hours_worked,
4                               na.rm = TRUE)) |>
5   ggplot(aes(x = avg_hours, y = education)) +
6   geom_bar(stat = "identity") +
7   scale_y_discrete(limits = c("Less than high
8                               school",
9                               "Some college",
10                              "Bachelor's degree",
11                              "Graduate degree"))
12   labs(title = "Working hours by education level",
13         x = "Average hours per week",
14         y = "Highest educational attainment")
```



Visualizing relationships

Two categorical variables: Faceted bar chart

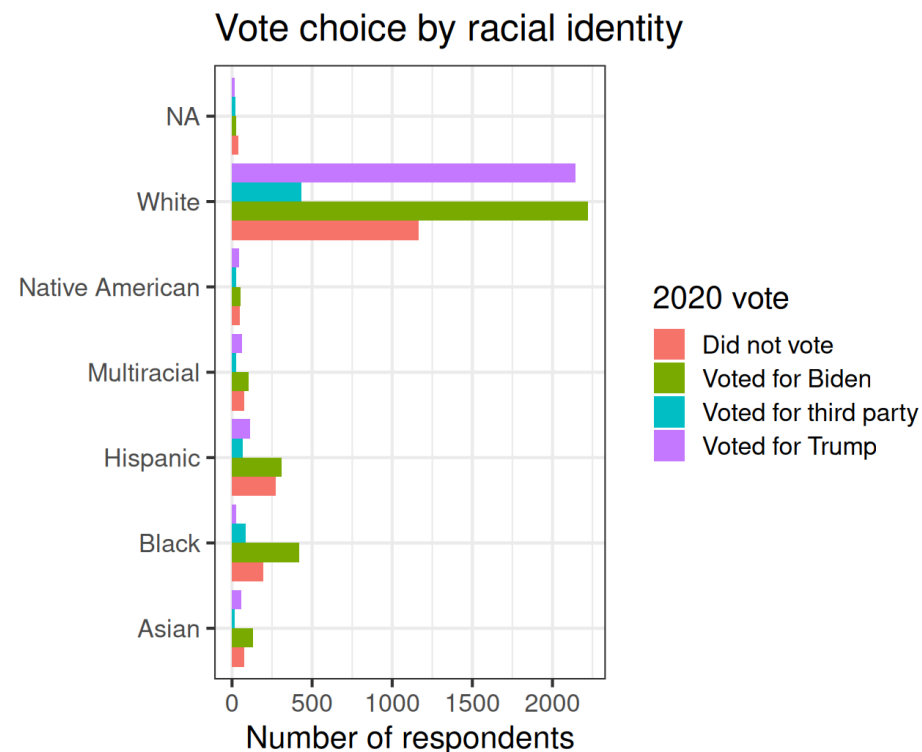
```
ggplot(df_anes, aes(y = vote_type)) +  
  geom_bar() +  
  facet_wrap(~ race, scales = "free_x") +  
  labs(x = "Number of respondents",  
       y = "",  
       title = "Vote choice by race in 2020 ANES")
```



Visualizing relationships

Two categorical variables: Dodged bar chart

```
ggplot(df_anes, aes(y = race)) +  
  geom_bar(aes(fill = vote_type),  
           position = "dodge") +  
  labs(x = "Number of respondents",  
       y = "",  
       fill = "2020 vote",  
       title = "Vote choice by racial identity")
```



Wrapping up

What we did today

1. Got data into R

- File directly on web \rightsquigarrow `read_csv("https://url.com/file.csv")`
- Otherwise \rightsquigarrow set working directory, save file there, `read_csv("file.csv")`

2. Manipulated data

- Subset by row with `filter()`, by column with `select()`
- Add or change columns with `mutate()`
- Calculate summaries with `group_by()` and `summarize()`
- Chain commands with the pipe `|>`

3. Visualized data with ggplot

- Histograms and bar charts for one-variable summaries
- Scatterplots, box plots, faceting for relationships

To do for next time

Next week's topic: **Causal questions and research design**

1. If anything from today was unfamiliar, practice with it
2. Read “Correlation, Causation, and Confusion” article
3. Read “Introduction to Causality” ebook chapter
4. Start thinking about topics you want to study in final project